

Alternate Derivation of Coalgebraic Representation of Büchi Automata Using Game Semantics

Research Internship Report

Jorrit de Boer

Abstract. We provide an explanation of existing literature for describing Büchi automata coalgebraically using trace semantics. To do this we also explain the modal μ -calculus and a coalgebraic model of nondeterministic systems. Finally, we present an alternate derivation of the coalgebraic model for Büchi automata using game semantics, which we believe is more intuitive than the one given in the original paper.

1 Introduction

Büchi automata and nondeterministic systems are crucial in theoretical computer science for modeling and verifying systems with infinite behaviors [3, 10]. Nondeterministic systems capture uncertainty and multiple outcomes, and are used in models like concurrent processes and nondeterministic Turing machines [7]. Büchi automata, which are in general also nondeterministic, handle infinite sequences of events, crucial for verifying systems that run indefinitely, such as operating systems or network protocols.

Coalgebra provides an effective framework for modeling state-based, dynamic systems. Techniques such as *coinduction* allow for reasoning about infinite structures, while *bisimulation* offers a formal way to establish behavioral equivalence between systems [8]. By modeling Büchi automata coalgebraically, these powerful tools can be applied for reasoning about infinite behaviors and nondeterminism.

The first goal of this report is to provide an understanding of the coalgebraic semantics using *trace semantics* of Büchi automata described in [9]. To do so we also explain the *modal μ -calculus*, a system for verifying properties of transition systems, and provide a coalgebraic model of nondeterministic systems, upon which the construction for the Büchi automata builds. By outlining these concepts we advance our first goal of the research internship, which is to gain an understanding of the current research into this topic.

Secondly we provide an alternate derivation of this coalgebraic representation using *game semantics*. Game semantics is a framework of describing a system in terms of a two-player game between a *verifier* and a *refuter* who want to verify, respectively refute, a statement [3]. By interpreting the modal μ calculus formulas which occur in the coalgebraic representation to a game we are able to use established theorems from game semantics to derive the coincidence between the coalgebraic model and the traces of the Büchi automata. We think that our approach provides a more intuitive proof of the results than the one provided in [9], which is quite cumbersome. Additionally, this formulation using game semantics might reveal connections to coalgebra automata which is based on game theoretic techniques [6].

The document is outlined as follows. In Section 2 we provide some background and relevant definitions for the rest of the report, which includes the modal μ -calculus and game semantics. In Section 3 we provide the coalgebraic representations of nondeterministic systems and Büchi automata from [4] and [9], respectively. In Section 4 we present our alternate derivation of the coincidence result given in the section before. Finally, in Section 5 we summarize the results and suggest directions for future work.

Contents

1 Introduction	1
2 Background	2
2.1 Büchi Automata	2
2.2 Parity Tree Automata	2
2.3 Fixed Points	3
2.4 Modal Mu-Calculus	4

2.5 Parity Games	5
3 Coalgebraic Representation of Büchi Automata	7
3.1 Finite Behavior of Nondeterministic Systems	7
3.2 Coalgebraic Representation of Büchi Automata	10
4 Derivation of Coincidence Using Game Semantics	11
5 Conclusion and Future Work	13
Bibliography	13
A Proofs	13

2 Background

2.1 Büchi Automata

Let us consider a very simple motivating example of a Büchi automaton, shown in Figure 1.




Figure 1: Example of a Büchi automaton.

This system represents some machine that takes requests, processes them, and returns some result. One might want to verify that this machine does not get stuck. In terms of the system shown, this would mean that the machine always ends up in the `idle` state again.

This behavior can be modeled using a Büchi automaton. A Büchi automaton, namely, is an automaton which models infinite behavior, and accepts those words for which there is a path through the automaton where the transitions are labeled by the letters of the word, and there is an accepting state that the path moves through infinitely many times. In this example, we make the `idle` state accepting, so the automaton accepts those words that always take the `return` transition again, and thus do not process indefinitely.

We can now give a formal definition of a Büchi automaton, and its *accepted language*:

Definition 2.1: A (nondeterministic) Büchi Automaton [3] is a tuple $A = \langle S, \Sigma, \delta, s_0, F \rangle$, with S a finite set of states, Σ the alphabet, $s_0 \in S$ the initial state, $\delta : S \times \Sigma \rightarrow \mathcal{P}(S)$ the transition function, $F \subseteq S$ the set of *final* (or *accepting*) states.

A *run* of a Büchi Automaton A on an ω -word $w = \sigma_0\sigma_1\ldots \in \Sigma^\omega$ is an infinite sequence of states $s_0, s_1, \ldots \in S^\omega$, such that s_0 is the initial state and for every $n \in \omega$, $s_{n+1} \in \delta(s_n, \sigma_n)$. A run is *accepting* if it passes through an accepting state infinitely many times. Equivalently (because F is finite), a run $\rho = s_0, s_1, \ldots$ is accepted if $\{i \mid s_i \in F\}$ is an infinite set. A word w is accepted by a Büchi automaton A if there is an accepting run of A on w . Finally, the accepted language $L(A)$ of a Büchi automaton, is the set of words accepted by A .

Indeed we now see that the accepted language for the example automaton is $(\text{request} \cdot \text{process}^* \cdot \text{return})^\omega$, where $*$ indicates repeating some set of letters/transitions some finite number of times (including zero) and ω indicates repeating indefinitely. That is, the machine gets a request, processes for at most some *finite* number of transitions and then returns some result. It does not get stuck processing indefinitely.

2.2 Parity Tree Automata

Büchi automata are actually a specific instance of parity tree automata. In this section we introduce this more general automaton. The coincidence results presented in Section 3.2 in fact not only hold for Büchi automata, but also for parity tree automata.

Instead of the acceptance criterion for Büchi automaton, we can use the parity acceptance condition. In this case, the states are not divided into accepting and non-accepting. Instead, every state has a priority, determined by $\Omega : S \rightarrow \omega$. A run $\rho = s_0, s_1, \ldots$ of an automaton A on a word w is then accepting if the maximum priority that occurs infinitely often is even. I.e., $\max\{\Omega(s) \mid s \text{ occurs infinitely often in } \rho\}$ is even. The Büchi acceptance criterion is the special case where non-accepting states have parity 1 and accepting states have parity 2.

Secondly, instead of words we can run our automaton on trees. In this case the alphabet Σ is *ranked* and has an arity function $|_|_ : \Sigma \rightarrow \omega$ indicating the number of branches a letter has. We denote the set of trees whose nodes are labeled with letters $\sigma \in \Sigma$ and whose branching is consistent with the arity of the letters as Tree_Σ . For example, if $|\sigma| = 2$ for all $\sigma \in \Sigma$, a tree $T \in \text{Tree}_\Sigma$ is binary tree with labels $\sigma \in \Sigma$. If $|\sigma| = 1$ for all $\sigma \in \Sigma$, Tree_Σ is just the set of infinite words over Σ .

We can now define a parity tree automaton:

Definition 2.2: A (nondeterministic) Parity Tree Automaton [3, 9] is a tuple $A = \langle S, \Sigma, \delta, s_0, \Omega \rangle$, with S a finite set of states, Σ a ranked alphabet with arity function $|_|_ : \Sigma \rightarrow \omega$, $s_0 \in S$ the initial state, $\delta : S \times \Sigma \rightarrow \mathcal{P}(S^*)$ the transition function where for each $\sigma \in \Sigma$ if $|\sigma| = n$ then $\delta(s)(\sigma) \subseteq S^n$, and $\Omega : S \rightarrow \omega$ that assigns a parity to each state.

A run ρ of the automaton A on a tree $T \in \text{Tree}_\Sigma$ is the tree T where the labels are replaced from letters $\sigma \in \Sigma$ to states $s \in S$ such that the root of the tree $\rho_0 = s_0$ is the initial state, and for a node in T with label $\sigma \in \Sigma$ the associated node in ρ with label $s \in S$ has children $s_1, \dots, s_{|\sigma|}$ such that $(s_1, \dots, s_{|\sigma|}) \in \delta(s)(\sigma)$. A run is accepted if for every branch of the tree, the maximum priority that occurs infinitely is even. A tree $T \in \text{Tree}_\Sigma$ is accepted by A if there is an accepting run of A on T . The accepted language of A is the set of accepted trees.

2.3 Fixed Points

Crucial for the next section, Section 2.4 about modal mu-calculus, is reasoning about *fixed points of monotone functions*. We briefly recall the important definitions and theorems.

Definition 2.3: A *complete lattice* is a partially ordered set $\langle L, \leq \rangle$ such that every subset $M \subseteq L$ has a least upper bound $\bigvee M$ and greatest lower bound $\bigwedge M$. Specifically, the whole set L has a least and greatest element, which we denote $\bigwedge L = \perp$ and $\bigvee L = \top$, respectively.

In this report we usually deal with the powerset of some set where subsets are ordered by inclusion. Indeed, for a set S , $\langle \mathcal{P}(S), \subseteq \rangle$ is a complete lattice. For $U \subseteq \mathcal{P}(S)$, $\bigvee U = \bigcup U$, and $\bigwedge U = \bigcap U$. The least and greatest elements are \emptyset and S , respectively.

Theorem 2.4 (Knaster-Tarski Fixed Point Theorem [1:Theorem 1.2.8]): Let $\langle L, \leq \rangle$ a complete lattice and $f : L \rightarrow L$ monotone ($f(x) \leq f(y)$ when $x \leq y$). Then, the set of fixed points $\{x \in L \mid f(x) = x\}$, is a complete lattice. Particularly, the function has a *least fixed point* (lfp) and a *greatest fixed point* (gfp).

There is a useful way of constructing these least and greatest fixed points. This is done by repeated function application on \perp for the least fixed point, and \top for the greatest fixed point. Concretely, we define for a monotone $f : L \rightarrow L$, for α an ordinal, and β a limit ordinal:

$$\begin{aligned} f^0 &:= \perp \\ f^{\alpha+1} &:= (f^\alpha) \\ f^\beta &:= \bigvee \{f^\alpha \mid \alpha < \beta\} \end{aligned} \tag{1}$$

This constructs an increasing chain

$$\perp = f^0 \leq f^1 \leq f^2 \leq \dots \tag{2}$$

which eventually stabilizes, giving the least fixed point, as stated by the following theorem:

Theorem 2.5 ([1:Theorem 1.2.11]): There exists an ordinal κ , such that $f^\kappa = f^{\kappa+1}$, which implies that f^κ is a fixed point of f . Furthermore, f^κ is the least fixed point of f . The dual process, beginning from \top and moving downward, constructs the greatest fixed point of f .

2.4 Modal Mu-Calculus

The modal mu-calculus is a powerful logic, used to verify properties of transition systems [1, 3]. We use it in Section 3.2 to select the right accepting trees for our coalgebraic system. In this section we give a concrete definition of modal mu-calculus formulas and provide intuition on how to use the modal mu-calculus to verify certain properties. We verify these properties over *transition systems*, which we define first:

Definition 2.6: A transition system (TS) is a tuple $T = \langle S, \delta, Prop, \lambda \rangle$ where S is the set of states, $\delta \subseteq S \times S$ the transition relation (we sometimes write $s \rightarrow s'$ if $(s, s') \in \delta$), $Prop$ the set of atomic propositions, and $\lambda : Prop \rightarrow \mathcal{P}(S)$ which interprets the atomic propositions.

You can see a TS as a directed graph where the vertices are labeled by atomic propositions $\mathcal{P}(Prop)$. Note that usually the modal mu-calculus is defined on *labeled* transition systems, but to simplify things slightly, and because we only need transition systems in the rest of the report we stick to transition systems.

Next we define the syntax of the modal mu-calculus:

Definition 2.7: A modal mu-calculus formula is defined by the grammar:

$$\varphi := P \mid \neg P \mid Z \mid \varphi_1 \wedge \varphi_2 \mid \varphi_1 \vee \varphi_2 \mid \Box \varphi \mid \Diamond \varphi \mid \mu Z. \varphi \mid \nu Z. \varphi \quad (3)$$

where $P \in Prop$ is an atomic proposition, $a \in \Sigma$ a label, and $Z \in Var$ a *fixed point variable*.

Note that you could define the modal mu-calculus without the \vee , $\langle a \rangle$, and ν operators, and define these instead in terms of the other operators, but we include them in the definition for legibility.

Definition 2.8: The semantics of a modal mu-calculus formula on a TS is a set of states where the formula holds, i.e. $|\varphi| \subseteq S$. For a modal mu-calculus formula φ , a transition system T , and an assignment $V : Var \rightarrow \mathcal{P}(S)$ we define:

$$\begin{aligned} \|P\|_V^T &:= \lambda(P) \\ \|\neg P\|_V^T &:= S \setminus \lambda(P) \\ \|Z\|_V^T &:= V(Z) \\ \|\varphi_1 \wedge \varphi_2\|_V^T &:= \|\varphi_1\|_V^T \cap \|\varphi_2\|_V^T \\ \|\varphi_1 \vee \varphi_2\|_V^T &:= \|\varphi_1\|_V^T \cup \|\varphi_2\|_V^T \\ \|\Box \varphi\|_V^T &:= \{s \mid \forall t \in S. \text{ if } s \rightarrow t \text{ then } t \in \|\varphi\|_V^T\} \\ \|\Diamond \varphi\|_V^T &:= \{s \mid \exists t \in S. \text{ if } s \rightarrow t \text{ then } t \in \|\varphi\|_V^T\} \\ \|\mu Z. \varphi\|_V^T &:= lfp(\lambda U. \|\varphi\|_{V[Z \mapsto U]}^T) = \bigcap \{U \subseteq X \mid U \subseteq \|\varphi\|_{V[Z \mapsto U]}^T\} \\ \|\nu Z. \varphi\|_V^T &:= gfp(\lambda U. \|\varphi\|_{V[Z \mapsto U]}^T) = \bigcup \{U \subseteq X \mid \|\varphi\|_{V[Z \mapsto U]}^T \subseteq U\} \end{aligned} \quad (4)$$

where $V[Z \mapsto U]$ is the valuation V except that Z maps to U .

We write $s \models^T \varphi$ if $s \in \|\varphi\|_V^T$ for an empty valuation V , or just $s \models \varphi$ if T is clear.

Let us briefly look at some intuition behind these definitions. We have $s \models^T p$ if in T at state s the propositional variable p holds. Conversely, $s \models^T \neg p$ holds if p does not hold in s . The \Diamond and \Box operators look at states reachable from s . For example, $s \models^T \Diamond p$ is true if there is some state s' such that $s \rightarrow s'$ and $s' \models^T p$. Analogously, $s \models^T \Box p$ is true if p is true in all successor states from s . Less intuitive are the μ and ν operators. Concretely, they identify least and greatest fixed points on functions from states to states. More intuitively, they can be used to define looping properties on transition systems, where μ can be used for finite looping, and ν for infinite looping. This will hopefully become more clear when looking at some examples:




Figure 2: Example of a TS. The sets next to the states denote the atomic propositions that hold in that state.

Consider the transition system given in Figure 2. We have $s_0 \models \Diamond p$, because there is a transition from s_0 to a state where p holds, namely $s_0 \rightarrow s_1$, because $s_1 \models p$. We, however, do not have $s_0 \models \Box p$, because $s_0 \rightarrow s_2$ and $s_2 \not\models p$.

To observe that μ is associated with finite looping, we look at the fact that $s_0 \models \mu Z.q \vee \Box Z$. This means that all finite paths from s_0 either reach a state with no outgoing transitions, or reach a state where q is true. We can see in Figure 2 that from s_0 every path reaches a state where q is true in finitely many steps. To more formally show that this holds, we make use of the method of constructing least and greatest fixed points in Theorem 2.5. The function we are calculating the lfp for is $f := \lambda U. \|q\| \cup \|\Box U\|$. The first iteration yields $f^1 = f(\emptyset) = \{s_2\}$, because $s_2 \models q$. Continuing, $f^2 = \{s_1, s_2\}$ and $f^3 = \{s_0, s_1, s_2\} = f^4$. So the lfp is the entire set of states S , and thus $s_0 \models \mu Z.q \vee \Box Z$.

Next we look at ν , which can be used for infinite looping. We show that $s_0 \not\models \nu Z.\Diamond Z$. This intuitively means that there exists an infinite path from s_0 . Indeed, we observe there are multiple infinite paths starting from s_0 . We confirm by computing the gfp: $f^1 = f(S) = \Diamond S = S$. Dually, observe that the lfp of this formula is $f^1(\emptyset) = \emptyset$. So we do not have $s_0 \models \mu Z.\Diamond Z$. This confirms the intuition that μ is for finite looping: there has to be some end point of the loop.

2.4.1 System of Equations

Next we introduce systems of equations with alternating fixed points. We only show how such a system works for two equations to save space and because that is all we use in the rest of the report. For more detail into this specific topic see [1, 9].

Definition 2.9: Let L_1, L_2 be partially ordered sets. An *equational system* is a system of two equations

$$u_1 \stackrel{=}{\eta_1} f_1(u_1, u_2) \quad u_2 \stackrel{=}{\eta_2} f_2(u_1, u_2) \quad (5)$$

where u_1, u_2 are variables, $\eta_1, \eta_2 \in \{\mu, \nu\}$, and $f_i : L_1 \times L_2 \rightarrow L_i$ are monotone functions. The solution to the system is defined by the following set of steps:

The intermediate solution $l_1^{(1)} := \eta_1 u_1.f_1(u_1, u_2)$, where we take the lfp if $\eta_1 = \mu$ and gfp if $\eta_1 = \nu$. Note that $l_1^{(1)} : L_2 \rightarrow L_1$.

The solution to the second equation is then given by $l^{\text{sol}} := \eta_2 u_2.f_2(l_1^{(1)}(u_2), u_2)$, where again we take the lfp if $\eta_2 = \mu$, and gfp if $\eta_2 = \nu$. The solution to the first equation is then $l_1^{\text{sol}} = l_1^{(1)}(l^{\text{sol}})$.

2.5 Parity Games

Next we introduce parity game and show how they can be used to give intuitive semantics for modal mu-calculus formulas. We use these semantics to prove the coincidence results in Section 4.

A parity game is a two player game between V (verifier) and R (refuter), who want to verify, respectively refute, a statement. In our case, this statement is $s \models^T \varphi$, i.e. that a modal mu-calculus formula holds in a state s in LTS T . So V wants to show $s \models^T \varphi$ and R wants to show $s \not\models^T \varphi$. The game consists of states and transitions between these states. Every state ‘belongs’ to either V or R , which determines what player picks the next transition is taken and thus the next state. A play of the game is then a (possibly infinite) sequence of states, and is won by either V or R . Concretely we define:

Definition 2.10 (Parity Game [3]): A parity game is a tuple $((S_V, S_R), E, \Omega)$, where $S = S_V \sqcup S_R$ is the set of states. From the states S_V player V picks the transition and for S_R player R does. $E \subseteq S \times$

S are transitions between the states. $\Omega : S \rightarrow \omega$ is the parity function, which determines the winner for infinite plays.

A play of the game is a (possibly infinite) sequence of states s_1, s_2, \dots such that $(s_i, s_{i+1}) \in E$. A finite play is won by a player if the other player gets stuck, i.e. has no moves from a position. An infinite play $\pi = s_1, s_2, \dots$ is won by V if $\max\{\Omega(s) \mid s \text{ occurs infinitely often in } \pi\}$ is even, and won by R if it is odd.

Next, we introduce the parity game for the modal mu-calculus. Consider the formula $\varphi = \varphi_1 \vee \varphi_2$. V wants to verify $s \models^T \varphi$, and to do so it suffices to show for either φ_i that $s \models^T \varphi_i$. Analogously for the formula $\varphi = \varphi_1 \wedge \varphi_2$, R can ‘pick’ the φ_i such that $s \not\models^T \varphi_i$, because if either φ_1 or φ_2 does not hold, φ does not hold. This same duality is seen in $\Diamond\varphi$ and $\Box\varphi$ where for $\Diamond V$ can show there is a transition for which φ holds, and for $\Box\varphi$, R can pick a transition such that φ does not hold. This way the game arises between V and R to determine whether $s \models^T \varphi_i$:

Definition 2.11 (Parity Game for Modal mu-Calculus[3]): For a transition system $T = (S, \rightarrow, \lambda)$ and a modal mu-calculus formula φ , we define the game $\mathcal{G}(\varphi, T) = ((S_V, S_R), E, \Omega)$ where:

- The states of the game $S_V \sqcup S_R = \{\varphi' \mid \varphi' \text{ is a subformula of } \varphi\} \times S$ are pairs of a subformula of φ and a state in the LTS. The subformula determines to what player the state belongs to. For a subformula ψ and a state s of the LTS:
 - $(\psi, s) \in S_V$ if
 - $\psi = \psi_1 \vee \psi_2$
 - $\psi = \Diamond\psi'$
 - $\psi = \eta Z.\psi'$ for $\eta \in \{\mu, \nu\}$
 - $\psi = Z$ for Z a fixed point variable
 - $\psi = p$ for p a propositional variable with $p \in \lambda(s)$.
 - $\psi = \neg p$ for p a propositional variable with $p \notin \lambda(s)$.
 - $(\psi, s) \in S_R$, if
 - $\psi = \psi_1 \wedge \psi_2$
 - $\psi = \Box\psi'$
 - $\psi = p$ for p a propositional variable with $p \notin \lambda(s)$.
 - $\psi = \neg p$ for p a propositional variable with $p \in \lambda(s)$.
- Edges E :
 - $(\psi_1 \vee \psi_2, s) \rightarrow (\psi_1, s)$ and $(\psi_1 \vee \psi_2, s) \rightarrow (\psi_2, s)$
 - $(\psi_1 \wedge \psi_2, s) \rightarrow (\psi_1, s)$ and $(\psi_1 \wedge \psi_2, s) \rightarrow (\psi_2, s)$
 - $(\Diamond\psi, s) \rightarrow (\psi, s')$ for any s' such that $s \rightarrow s'$ in T .
 - $(\Box\psi, s) \rightarrow (\psi, s')$ for any s' such that $s \rightarrow s'$ in T .
 - $(\eta Z.\psi, s) \rightarrow (\psi, s)$ and $(Z, s) \rightarrow (\psi, s)$ for $\eta \in \{\mu, \nu\}$
- The priority function Ω depends on the *alternation depth* $\alpha(\psi)$ of the subformula ψ , which is defined as follows:
 - $\alpha(p) = \alpha(\neg p) = 0$ for p a propositional variable
 - $\alpha(\psi_1 \wedge \psi_2) = \alpha(\psi_1 \vee \psi_2) = \max\{\alpha(\psi_1), \alpha(\psi_2)\}$
 - $\alpha(\Diamond\psi) = \alpha(\Box\psi) = \alpha(\psi)$
 - $\alpha(\mu Z.\psi) = \max(\{1, \alpha(\psi)\} \cup \{\alpha(\nu Z'.\psi' + 1) \mid \nu Z'.\psi' \text{ is a subformula of } \psi \text{ and } Z \text{ occurs free in } \psi'\})$
 - $\alpha(\nu Z.\psi) = \max(\{1, \alpha(\psi)\} \cup \{\alpha(\mu Z'.\psi' + 1) \mid \mu Z'.\psi' \text{ is a subformula of } \psi \text{ and } Z \text{ occurs free in } \psi'\})$

Intuitively, the alternation depth of a formula is the maximum number of alternating μ/ν operators, where we only count those alternations where the free variable actually occurs freely in the subformula, meaning the fixed point operators are actually interdependent. Ω is then:

- $\Omega((\mu Z.\psi, s)) =$ the smallest odd number greater or equal than $\alpha(\psi) - 1$
- $\Omega((\nu Z.\psi, s)) =$ the smallest even number greater or equal than $\alpha(\psi) - 1$
- $\Omega((\psi, s)) = 0$ iff ψ is not a μ or ν formula.

Where the intuition for operators like $\vee, \wedge, \square, \Diamond$ is quite straightforward, for the μ/ν operators it is less so. Briefly put, it follows from what was explained in Section 2.4 that μ incites finite looping, and ν infinite looping. It can be seen from the definition for Ω using the alternation depth, that outer μ/ν operators have higher priority than inner ones, and μ is always even and ν odd. Thus the highest priority occurring infinitely often in an infinite play indicates the outermost fixed point operator that is visited infinitely often. Thus, if this is even, we have an infinite loop through a ν operator, which satisfies the formula. For a μ operator, however, an infinite loop is undesired, and thus if the outermost fixed point operator which is visited infinitely often is μ , it is not a least fixed point, and R has refuted the formula.

Now, to use this game to give alternative semantics for the modal mu-calculus we need that if $s \models^T \varphi$ then V can verify this in the game $\mathcal{G}(\varphi, T)$ by winning the game, and R can not win. We call this that V has a winning strategy: V can always play (i.e. take the right transition if it is their turn) such that regardless of what R plays, V wins the play. We then have the theorem, which is crucial for our derivation of the coincidence results in Section 4:

Theorem 2.12 (Theorem 10.18 [3]):

$$s \models^T \varphi \Leftrightarrow V \text{ has a winning strategy in } \mathcal{G}(\varphi, T) \text{ starting in state } (\varphi, s) \quad (6)$$

The thorough explanation and rigorous proof are quite intricate, so to keep the presentation simple we limit ourselves to this intuition given above.

3 Coalgebraic Representation of Büchi Automata

3.1 Finite Behavior of Nondeterministic Systems

In this section we present a coalgebraic representation of nondeterministic systems. The next section for Büchi automata builds upon this construction.

3.1.1 Deterministic Automata

First we consider a deterministic finite automaton, $\langle S, \Sigma, \delta, o \rangle$ with S the states, Σ the alphabet, $\delta : S \times \Sigma \rightarrow S$ the transition function, and $o : S \rightarrow 2$ with $2 = \{0, 1\}$, the output function determining if a state is final. We do not consider an initial state here because we just want to obtain the accepted words for each state. Such an automaton can be represented by a coalgebra $c : S \rightarrow 2 \times S^\Sigma$ for the functor $F(S) = 2 \times S^\Sigma$. This is a very useful construction because a final coalgebra for this functor is carried by 2^{Σ^*} , and the unique coalgebra homomorphism beh to this final coalgebra captures exactly the language accepted by a state [8]. This is shown in the commuting diagram:

$$\begin{array}{ccc} 2 \times S^\Sigma & \xrightarrow{\text{id} \times \text{beh}^\Sigma} & 2 \times (2^{\Sigma^*})^\Sigma \\ \uparrow \langle o, \delta \rangle & & \uparrow \langle e, d \rangle \\ S & \xrightarrow{\text{beh}} & 2^{\Sigma^*} \end{array} \quad (7)$$

Here $e : 2^{\Sigma^*} \rightarrow 2$ is given by $e(L) = L(\varepsilon)$, i.e., $e(L) = 1$ iff L contains the empty word. And $d : 2^{\Sigma^*} \rightarrow (2^{\Sigma^*})^\Sigma$ is the language derivative, given by $d(L)(a) = L_a$ where $L_a(w) = L(aw)$, so $w \in d(L)(a) = L_a$ iff $aw \in L$.

Working out the paths through the diagram we obtain that

- $\text{beh}(s)(\varepsilon) = o(s)$, and
- $\text{beh}(s)(\sigma w) = \text{beh}(\delta(s)(\sigma))(w)$,

for $s \in S$, $\sigma \in \Sigma$, $w \in \Sigma^*$. So $\text{beh}(s)$ contains the empty word iff s is a final state, and accepts σw iff $\delta(s)(\sigma)$ accepts w . Which is precisely the language accepted by state x in the deterministic finite automaton!

3.1.2 Finite Behavior Nondeterministic Automata

Unfortunately, extending this approach to nondeterministic systems is not possible, as we will illustrate by the following system, which we will use as a running example:




Figure 3: Example of a nondeterministic automaton.

The automaton given in Figure 3 is nondeterministic because in s_0 there are two transitions for a . Intuitively, the finite words accepted by the system from state s_0 should be

$$\text{beh}(s_0) = \{a, ab, abb, \dots\} \cup \{a, ac, acc, \dots\}. \quad (8)$$

This transitions system might be modeled by a coalgebra $c : S \rightarrow 2 \times \mathcal{P}(\Sigma \times S)$, i.e., for every state whether it is final, and a set of pairs $(\sigma, s) \in \Sigma \times S$ denoting a transition by taking letter σ and transitioning to state s . The problem is that this functor $FS = 2 \times \mathcal{P}(\Sigma \times S)$ does not have a final coalgebra, as Lambek's lemma implies that such a final coalgebra $z : Z \rightarrow 2 \times \mathcal{P}(\Sigma \times Z)$ for some carrier Z , would have to be an isomorphism [2]. But an isomorphism $Z \cong 2 \times \mathcal{P}(\Sigma \times Z)$ would imply a bijection between Z and $\mathcal{P}(Z)$, which cannot exist.

The solution, as given by Hasuo et al. [4], is to work in the Kleisli category for the monad \mathcal{P} . Recall that a map $f : X \rightarrow Y$ in the Kleisli category is a map $f : X \rightarrow \mathcal{P}(Y)$ in **Sets**. Briefly put, this will solve our problem because we can have a final coalgebra $z : Z \rightarrow FZ$ that is a map $z : Z \rightarrow \mathcal{P}(FZ)$ in **Sets**. Next, we will review the definition of the Kleisli category and define the appropriate functor, enabling us to construct the desired final coalgebra that characterizes the accepting finite words.

The powerset monad \mathcal{P} is defined by the unit $\eta_X : X \rightarrow \mathcal{P}(X)$ which sends an element of X to the singleton set, $\eta_X(x) = \{x\}$ for $x \in X$, and the multiplication $\mu_X : \mathcal{P}(\mathcal{P}(X)) \rightarrow \mathcal{P}(X)$ which takes the union of the sets, i.e. $\mu_X(A) = \bigcup_{a \in A} a$. For a function $f : X \rightarrow Y$ we get $\mathcal{P}(f) : \mathcal{P}(X) \rightarrow \mathcal{P}(Y)$ by $\mathcal{P}(f)(A) = \{f(a) \mid a \in A\}$. The Kleisli category for this monad is defined as follows:

- **objects:** the same as for **Sets**, sets
- **morphisms:** a morphism f from X to Y in $\mathcal{Kl}(\mathcal{P})$ is a map $f : X \rightarrow \mathcal{P}(Y)$ in **Sets**. For morphisms $f : X \rightarrow Y$ and $g : Y \rightarrow Z$ in $\mathcal{Kl}(\mathcal{P})$ (so $f : X \rightarrow \mathcal{P}(Y)$ and $g : Y \rightarrow \mathcal{P}(Z)$ in **Sets**) we define $(g \circ f)$ in $\mathcal{Kl}(\mathcal{P})$ as $(\mu_Z \circ \mathcal{P}(g) \circ f) : X \rightarrow \mathcal{P}(Z)$, so $(g \circ f) : X \rightarrow Z$ in $\mathcal{Kl}(\mathcal{P})$.

Next, we construct our functor in $\mathcal{Kl}(\mathcal{P})$, which we call the lifting of F in $\mathcal{Kl}(\mathcal{P})$, and denote \overline{F} . The key here is that because we are working in the Kleisli category, if we use the functor $\overline{F}S = \Sigma \times S$, the coalgebra map $c : S \rightarrow \Sigma \times S$, will be a map $c : S \rightarrow \mathcal{P}(\Sigma \times S)$ in **Sets**, which models nondeterministic transitions. In the previous section we used the functor $FS \rightarrow 2 \times S^\Sigma$ where $o : S \rightarrow 2$ denoted whether the state was final. Combining this with the functor $\overline{F}S = \Sigma \times S$ in the Kleisli category we would get the functor $\overline{F}S = 2 \times \Sigma \times S$ and the coalgebra $c : S \rightarrow 2 \times \Sigma \times S$ which is $c : S \rightarrow \mathcal{P}(2 \times \Sigma \times S)$ in **Sets**, which would mean that every transition can be final or not, which is not what we want. For this reason we use the functor $\overline{F}S = 1 + \Sigma \times S$ such that the coalgebra $c : S \rightarrow 1 + \Sigma \times S$ is the map $c : S \rightarrow \mathcal{P}(1 + \Sigma \times S)$ where $s \in S$ is final iff $* \in c(s)$ (note that we use $1 = \{*\}$).

This works easily on objects, $\overline{F}X = FX$, because in the Kleisli category, the objects are the same. But for morphisms we have to do a little bit more work. Observe that because a map $f : X \rightarrow Y$ in $\mathcal{Kl}(\mathcal{P})$ is a map $f : X \rightarrow \mathcal{P}(Y)$ in **Sets**, applying the functor F on the map itself would yield $Ff : FX \rightarrow F\mathcal{P}(Y)$. So what we need is a natural transformation $\lambda : F\mathcal{P} \Rightarrow \mathcal{P}F$, i.e., a distributive law [4], such that $1 + \Sigma \times (\mathcal{P}(S)) \xrightarrow{\lambda} \mathcal{P}(1 + \Sigma \times X)$. We define this as $* \mapsto \{*\}$, and $(\sigma, S) = \{(\sigma, x) \mid x \in S\}$ for $\sigma \in \Sigma$ and $S \subseteq X$. This follows intuitively if you observe that if from state s taking transition σ takes you to $\{x, y, z\}$ ($(\sigma, \{x, y, z\}) \in c(s)$, or $\{x, y, z\} \in \delta(s)(\sigma)$), you can also see this as transitions $\{(\sigma, x), (\sigma, y), (\sigma, z)\}$.

Finally, the main theorem from [4] (Theorem 3.3), and the last ingredient to make the construction work is that the initial algebra for the functor F in **Sets**, gives us the final coalgebra for the lifted functor \overline{F} in $\mathcal{Kl}(\mathcal{P})$. Specifically, for this functor $FS = 1 + \Sigma \times S$ and its lifting as described above, the initial F -algebra $\alpha : FA \rightarrow A$ in **Sets** yields a final \overline{F} -coalgebra in $\mathcal{Kl}(\mathcal{P})$ by:

$$(\eta_{FA}\alpha)^{-1} = \eta_{FA}\alpha^{-1} : A \rightarrow \overline{F}A \text{ in } \mathcal{Kl}(\mathcal{P}) \quad (9)$$

In fact, this result holds more generally: for the lifting monad \mathcal{L} , the subdistribution monad \mathcal{D} , and any shapely functor F , see [4] for more details.

The initial F -algebra for our functor $FS = 1 + \Sigma \times S$ in **Sets** is $[\text{nil}, \text{cons}] : 1 + \Sigma \times \Sigma^* \rightarrow \Sigma^*$. So we get the commuting diagram

$$\begin{array}{ccc}
 1 + \Sigma \times S & \xrightarrow{1 + \Sigma \times \text{beh}} & 1 + \Sigma \times \Sigma^* \\
 \uparrow c & & \uparrow \cong \\
 S & \xrightarrow{\text{beh}} & \Sigma^*
 \end{array}
 \quad \eta_{1+\Sigma \times \Sigma^*} \circ [\text{nil}, \text{cons}]^{-1} \quad \text{in } \mathcal{KL}(\mathcal{P}).
 \tag{10}$$

Following the paths within the diagram we obtain that

$$\begin{aligned}
 \varepsilon \in \text{beh}(s) &\iff * \in c(s) \iff \text{state } s \text{ is accepting} \\
 \sigma w \in \text{beh}(s) &\iff \exists t. (s \xrightarrow{\sigma} t \wedge w \in \text{beh}(t)).
 \end{aligned}
 \tag{11}$$

Explained in words, a state accepts the empty word iff the state is accepting, and it accepts σw for $\sigma \in \Sigma$ and $w \in \Sigma^*$ iff it can transition with σ to a state which accepts w . Which is exactly the desired words!

3.1.3 Possibly Infinite Behavior

As a step towards infinite words in Büchi automata let us consider infinite words in Figure 3. We can slightly alter our previous construction to additionally obtain infinite words through this system. Concretely, the infinite words for the system in Figure 3 for x_0 are ab^ω and ac^ω .

The intuition for this new construction is as follows. In the previous section we constructed the final coalgebra for the lifted functor \overline{F} using the initial F -algebra in **Sets**. In the example of the LTS with termination the initial algebra was carried by Σ^* . The final coalgebra in **Sets** for F is carried by Σ^∞ (where the ∞ operators means some finite number of times or indefinitely so $\Sigma^\infty = \Sigma^* \cup \Sigma^\omega$) the set of finite and infinite words. So if we use this final coalgebra instead of the initial algebra, do we obtain both the finite and infinite words?

Consider again the monad \mathcal{P} , our functor F (this too holds more generally, see [4, 5]), and its lifting in the Kleisli category \overline{F} . For a final coalgebra $\xi : Z \rightarrow FZ$, the coalgebra

$$\eta_{FZ} \circ \xi : Z \rightarrow \overline{F}Z \text{ in } \mathcal{KL}(\mathcal{P})
 \tag{12}$$

is *weakly final*. That means, for any coalgebra $c : S \rightarrow \overline{F}S$, there is a morphism $\text{beh} : S \rightarrow Z$ in $\mathcal{KL}(\mathcal{P})$ such that the following diagram commutes

$$\begin{array}{ccc}
 \overline{F}S & \xrightarrow{\overline{F}(\text{beh})} & \overline{F}Z \\
 \uparrow c & & \uparrow \cong \\
 S & \xrightarrow{\text{beh}} & Z
 \end{array}
 \quad \eta_{FZ} \circ \xi \quad \text{in } \mathcal{KL}(\mathcal{P}),
 \tag{13}$$

but this morphism is not necessarily unique. However, there is a canonical choice beh^∞ among these morphisms, namely the one which is maximal with respect to inclusion. We call this function $\text{beh}^\infty : S \rightarrow \mathcal{P}(Z)$ (in **Sets**) the *possibly-infinite* behavior for c .

Indeed, if we consider our running example Figure 3 with termination, $\xi : \Sigma^\infty \rightarrow 1 + \Sigma \times \Sigma^\infty$ is the final F -coalgebra, defined by $\xi(\varepsilon) = *$ and $\xi(\sigma w) = (\sigma, w)$. Instantiating the diagram in Equation 13, we obtain

$$\begin{aligned}
 \varepsilon \in \text{beh}^\infty(s) &\iff * \in c(s) \iff \text{state } s \text{ is accepting} \\
 \sigma w \in \text{beh}^\infty(s) &\iff \exists t. (s \xrightarrow{\sigma} t \wedge w \in \text{beh}^\infty(t)).
 \end{aligned}
 \tag{14}$$

Which is the same as in Equation 11. However, because the domain is Σ^∞ , we obtain different words when we take the maximal function satisfying these equations. Namely the finite words, in addition to the infinite ones! For the system in Figure 3 we get the same words as before, but additionally $\{ab^\infty, ac^\infty\} \subseteq \text{beh}_c^\infty(s_0)$. Interestingly, taking the minimum morphism we again obtain just the finite words [4, 5].

3.2 Coalgebraic Representation of Büchi Automata

We can apply the previous framework for possibly infinite words to our initial example for a Büchi automaton, in Figure 1. This would yield all infinite words through the automaton, so also for example $\text{request} \cdot \text{process}^\omega$, it does not take into account accepting states, only for ending finite words. How do we eliminate those words that process indefinitely? That is, only accept those words under the Büchi acceptance criterion of passing through an accepting state infinitely many times.

A way of solving this is given by [9]. In short, the main idea of this paper is to divide the states into accepting and non-accepting states. Then, applying the previous construction using the final F -coalgebra in **Sets** we obtain two separate commuting diagrams for these disjoint sets of states. And finally, using greatest and least fixed points we can precisely pick exactly the accepting words for the Büchi automaton.

We first give the commuting diagrams which govern the behavior mappings. We are now considering Büchi automata, so the functor we consider is $FS = \Sigma \times S$, the final coalgebra for this functor is $d : \Sigma^\omega \rightarrow \Sigma \times \Sigma^\omega$, defined by $d(\sigma \cdot w) = (\sigma, w)$, and the monad is still \mathcal{P} . The lifting \bar{F} is effectively the same, just without a case for $*$ $\in 1$. We now consider the state space as a disjoint union $S = S_1 \cup S_2$ of non-accepting and accepting states, respectively. This gives rise to two separate coalgebras $c_i : S_i \rightarrow \bar{F}X$, defined by the restriction $c \circ \kappa_i : S_i \rightarrow \bar{F}X$ along the coprojection $\kappa_i : S_i \hookrightarrow S$ for $i \in \{1, 2\}$. We then get the two commuting diagrams:

$$\begin{array}{ccc}
 \Sigma \times [\text{beh}_1, \text{beh}_2] & & \Sigma \times [\text{beh}_1, \text{beh}_2] \\
 \Sigma \times S \xrightarrow{\quad} \Sigma \times \Sigma^\omega & & \Sigma \times S \xrightarrow{\quad} \Sigma \times \Sigma^\omega \\
 \uparrow c_1 & \begin{array}{c} \cong \\ \mu \end{array} & \uparrow \eta_{\Sigma^\omega} \circ d \\
 S_1 \xrightarrow[\text{beh}_1]{\quad} \Sigma^\omega & & S_2 \xrightarrow[\text{beh}_2]{\quad} \Sigma^\omega
 \end{array}
 \quad \text{in } \mathcal{K}\ell(\mathcal{P}). \tag{15}$$

Where $\overset{\mu}{=}$ and $\overset{\nu}{=}$ mean that we take the least behavior mapping in the left diagram to obtain beh_1 , and the greatest behavior mapping in the right diagram to obtain beh_2 . More concretely, $\text{beh}_1 : S_1 \rightarrow \mathcal{P}(\Sigma^\omega)$ and $\text{beh}_2 : S_2 \rightarrow \mathcal{P}(\Sigma^\omega)$, are the solutions to the following system of equations:

$$\begin{aligned}
 u_1 &\overset{\mu}{=} (\eta_{\Sigma^\omega} \circ d)^{-1} \odot \bar{F}[u_1, u_2] \odot c_1 \\
 u_2 &\overset{\nu}{=} (\eta_{\Sigma^\omega} \circ d)^{-1} \odot \bar{F}[u_1, u_2] \odot c_2
 \end{aligned} \tag{16}$$

We first rewrite this to something more clear and usable:

Lemma 3.1: The traces in Equation 16 coincide with:

$$u_1 \overset{\mu}{=} \Diamond_\delta([u_1, u_2]) \upharpoonright S_1 \quad u_2 \overset{\nu}{=} \Diamond_\delta([u_1, u_2]) \upharpoonright S_2 \tag{17}$$

Where $\Diamond_\delta : (\mathcal{P}(\Sigma^\omega))^S \rightarrow (\mathcal{P}(\Sigma^\omega))^S$ is given by

$$\Diamond_\delta(\text{beh})(s) = \{\sigma \cdot w \mid s' \in \delta(s)(\sigma), w \in \text{beh}(s')\}. \tag{18}$$

By taking exactly those behavior mappings which are the solution to this system of equation, we take exactly those words that the Büchi automaton accepts:

Lemma 3.2 ([9:Lemma 4.5]): Let $A = (S, \Sigma, \delta, s_0, F)$ be a Büchi automaton, where we let $S = S_1 \cup S_2$ the disjunct union of the non-accepting and accepting states, respectively, so $S_1 = S \setminus F$, $S_2 = F$. Let $l_1^{\text{sol}}, l_2^{\text{sol}}$ be the solutions to the following equational system, where the variables u_1, u_2 range over $(\mathcal{P}(\Sigma^\omega))^{S_i}$

$$u_1 \stackrel{\mu}{=} \Diamond_\delta^1(u_1 \cup u_2) \quad u_2 \stackrel{\nu}{=} \Diamond_\delta^2(u_1 \cup u_2) \quad (19)$$

Where $\Diamond_\delta^i : (\mathcal{P}(\Sigma^\omega))^{S_i} \rightarrow (\mathcal{P}(\Sigma^\omega))^{S_i}$ is given by

$$\Diamond_\delta^i(\text{beh})(s) = \{\sigma \cdot w \mid \sigma \in \Sigma, s' \in \delta(s)(\sigma), w \in \text{beh}(s')\}. \quad (20)$$

Then the solutions $l_i^{\text{sol}} : S_i \rightarrow \Sigma^\omega$ map S_i to the accepted language from that state.

We provide a brief intuition here, utilizing what was observed in Section 2.4. Namely, that μ is associated with finite looping, and ν with infinite. So the second equation makes sure the run passes through S_2 infinitely many times. Note that it can still move through S_1 , but it has to move through S_2 infinitely many times. The first equation, with the μ operator, makes sure that any run passing through S_1 passes to the second equation in some finite number of steps, where it passes through S_2 infinitely many times. So the two equations make sure that a run passes through S_2 (the second equation) infinitely many times, and when it passes through S_1 it passes back to S_2 in a finite number of steps where it can pass through S_2 infinitely many times again.

Regardless of this intuition, the proof of this lemma given in [9] is rather complex. In the next section we provide our proof using game semantics, which we believe is a lot more comprehensive.

Combining Lemma 3.1 and Lemma 3.2 we obtain the coincidence result:

Theorem 3.3 ([9:Theorem 4.6]): Let $A = (S, \Sigma, \delta, s_0, F)$ be a Büchi automaton. Then the behavior mappings $\text{beh}_1, \text{beh}_2$, which are the solution to the system of equations in Equation 16 coincide with the accepted language of A : $\text{beh}(s_0) = [\text{beh}_1, \text{beh}_2](s_0) = L(A)$.

4 Derivation of Coincidence Using Game Semantics

In this section we prove Theorem 3.3 using game semantics and in a very pretty way.

Definition 4.1: Let $A = (X_1 \cup X_2, \Sigma, \delta)$ be a Büchi automaton, with states $X = X_1 \cup X_2$ where X_2 are the accepting states, Σ is the alphabet, and $\delta : X \times \Sigma \rightarrow \mathcal{P}(X)$ the transition function. We define a Transition System (TS) over the set of propositional variables $\{p_1, p_2\}$ for this automaton, denoted as T_A , as follows:

- States are (x, w) for $x \in X$ and $w \in \Sigma^\omega$
- Transition $(x, \sigma w) \rightarrow (x', w)$ for $x, x' \in X, \sigma \in \Sigma, w \in \Sigma^\omega$, iff $x' \in \delta(x)(\sigma)$
- Labeling function given by $\lambda((x, w)) = \{p_i\}$ iff $x \in X_i$, i.e. the propositional variables denote for what i , we have $x \in X_i$.

We observe how the formulas from Lemma 3.1 are built up and convert them. For example the closed formula for $\text{beh}_1 : X_1 \rightarrow \Sigma^\omega$ for $n = 2$ priorities, i.e. a Büchi automaton, $\text{beh}_1 = \nu u_2. \Diamond_\delta[\mu u_1. [u_1, u_2] \uparrow X_1, u_2] \uparrow X_2$. Let φ be such a formula, then:

- $\varphi = u$ a free variable, or
- $\varphi = \Diamond_\delta \varphi'$, or
- $\varphi = \eta u. \varphi'$ where $\eta \in \{\mu, \nu\}$, or
- $\varphi = \varphi' \uparrow X_i$, or
- $\varphi = [\varphi_1, \dots, \varphi_n]$

we also observe the implicit semantics of the formula $\|\varphi\|$: ...

Definition 4.2: So we convert a formula φ , to our desired formula $\bar{\varphi}$ to conform to Definition 10.2[]:

- $\varphi = u$ a free variable then $\bar{\varphi} = u$ also a free variable
- $\varphi = \Diamond_\delta \varphi'$ then $\bar{\varphi} = \Diamond \bar{\varphi}'$
- $\varphi = \eta u. \varphi'$ for $\eta \in \{\mu, \nu\}$ then $\bar{\varphi} = \eta u. \bar{\varphi}'$

- $\varphi = \varphi' \uparrow X_i$ then $\overline{\varphi} = p_i \wedge \overline{\varphi'}$
- $\varphi = [\varphi_1, \dots, \varphi_n]$ then $\overline{\varphi} = (p_1 \wedge \overline{\varphi_1}) \vee \dots \vee (p_n \wedge \overline{\varphi_n})$

Lemma 4.3: For a modal μ -calculus formula φ (a la paper 1) and a valuation $V : Var \rightarrow (X \rightarrow \mathcal{P}(\Sigma^\omega))$, $x \in X, w \in \Sigma^\omega$:

$$w \in \|\varphi\|_V(x) \Leftrightarrow (x, w) \in \|\overline{\varphi}\|_{\overline{V}}^{TA} \quad (21)$$

where $\overline{V}(U) = \{(x, w) \mid x \in X, w \in V(U)(x)\}$

Proof: We prove this by induction on the formula φ . The base case is $\varphi = U$ a free variable:

$$w \in \|U\|_V(x) = V(U)(x) \Leftrightarrow (x, w) \in \overline{V}(U) = \|U\|_{\overline{V}}^{TA}$$

Induction step:

- $\varphi = \mu U. \varphi'$:

We have to show $w \in \|\mu U. \varphi'\|_V(x) = \text{lfp}(\lambda u. \|\varphi'\|_{V[U \mapsto u]}) \Leftrightarrow (x, w) \in \|\mu U. \overline{\varphi'}\|_{\overline{V}} = \text{lfp}(\lambda u. \|\overline{\varphi'}\|_{\overline{V}[U \mapsto u]})$. Let $W = \text{lfp}(\lambda u. \|\varphi'\|_{V[U \mapsto u]})$. We define $\overline{W} = \{(x, w) \mid x \in X, w \in W(x)\}$ and show $W = \text{lfp}(\lambda u. \|\varphi'\|_{V[U \mapsto u]}) \Leftrightarrow \overline{W} = \text{lfp}(\lambda u. \|\overline{\varphi'}\|_{\overline{V}[U \mapsto u]})$. For this we first prove that W is a fixed point iff \overline{W} is a fixed point:

Assume W is a fixed point, so $\|\varphi'\|_{V[U \mapsto W]} = W$. We observe that for a valuation V and V' where $V' = V[U \mapsto W]$, we have the converted valuation $\overline{V'} = \overline{V}[U \mapsto \overline{W}]$. We use this to incite the IH to get $w \in \|\varphi'\|_{V[U \mapsto W]} \Leftrightarrow (x, w) \in \|\overline{\varphi'}\|_{\overline{V}[U \mapsto \overline{W}]}$. Using this we get $(x, w) \in \|\overline{\varphi'}\|_{\overline{V}[U \mapsto \overline{W}]} \Leftrightarrow w \in \|\varphi'\|_{V[U \mapsto W]}(x) = W(x) \Leftrightarrow (x, w) \in \overline{W}$, so $\|\overline{\varphi'}\|_{\overline{V}[U \mapsto \overline{W}]} = \overline{W}$, so \overline{W} is a fixed point.

Now assume \overline{W} is a fixed point, so $\|\overline{\varphi'}\|_{\overline{V}[U \mapsto \overline{W}]} = \overline{W}$. Then, for $x \in X, W(x) = \{w \mid (x, w) \in \overline{W}\}$. Applying IH like the previous case again we obtain $w \in \|\varphi'\|_{V[U \mapsto W]}(x) \Leftrightarrow (x, w) \in \|\overline{\varphi'}\|_{\overline{V}[U \mapsto \overline{W}]} = \overline{W} \Leftrightarrow w \in W(x)$. So $w \in \|\varphi'\|_{V[U \mapsto W]}(x) \Leftrightarrow w \in W(x)$ for all $x \in X$, so $\|\varphi'\|_{V[U \mapsto W]} = W$, so W is a fixed point.

Next, we show that W is the *least* fixed point iff \overline{W} is the *least* fixed point:

Assume W is a lfp, from above we know that \overline{W} is a fixed point. Take some other fixed point \overline{Y} , i.e. $\|\overline{\varphi'}\|_{\overline{V}[U \mapsto \overline{Y}]} = \overline{Y}$. Now, again inciting what we showed above, we know Y is a fixed point, so $\|\varphi'\|_{V[U \mapsto Y]} = Y$. So because W is the lfp, for all $x, W(x) \subseteq Y(x)$. From this it follows that $(x, w) \in \overline{W} \rightarrow w \in W(x) \rightarrow w \in Y(x) \rightarrow (x, w) \in \overline{Y}$, so $\overline{W} \subseteq \overline{Y}$. So \overline{W} is the least fixed point.

For the other direction, assume \overline{W} is a least fixed point. Then W is a fixed point. Take some other fixed point Y , i.e. $\|\varphi'\|_{V[U \mapsto Y]} = Y$, then \overline{Y} is a fixed point. So because \overline{W} is the lfp, we have $\overline{W} \subseteq \overline{Y}$. Now for any w, x we have $w \in W(x) \rightarrow (x, w) \in \overline{W} \rightarrow (x, w) \in \overline{Y} \rightarrow w \in Y(x)$. So $W \subseteq Y$.

- $\varphi = \nu U. \varphi'$:

This case is a analogous to the μ case. The first part proving W is a fixed point iff \overline{W} is a fixed point, and for proving W is a *greatest* fixed point iff \overline{W} is too you reason in the opposite direction as for μ .

- $\varphi = \Diamond_\delta \varphi'$:

$$w \in \|\Diamond_\delta \varphi'\|_V(x) = \{\sigma w \mid \exists x' \in \delta(x)(\sigma)[w \in \|\varphi'\|_V(x')]\} \stackrel{IH}{=} \{\sigma w \mid \exists x' \in \delta(x)(\sigma)[(x', w) \in \|\overline{\varphi'}\|_{\overline{V}}]\} \Leftrightarrow (x, w) \in \{(x, \sigma w) \mid \exists x' \in \delta(x)(\sigma)[(x', w) \in \|\overline{\varphi'}\|_{\overline{V}}]\} = \|\Diamond_\delta \overline{\varphi'}\|_{\overline{V}}^{TA}$$

- $\varphi = \varphi' \upharpoonright X_i$:

$$w \in \|\varphi' \upharpoonright X_i\|_V(x) \Leftrightarrow x \in X_i \wedge w \in \|\varphi'\|_V(x) \stackrel{IH}{\Leftrightarrow} x \in X_i \wedge (x, w) \in \|\overline{\varphi'}\|_{\overline{V}} \Leftrightarrow (x, w) \in \|p_i \wedge \overline{\varphi'}\|_{\overline{V}}$$

- $\varphi = [\varphi_1, \dots, \varphi_n]$:

$\|\varphi\|_V(x) = \begin{cases} \|\varphi_1\|_V(x) & \text{if } x \in X_1 \\ \vdots \\ \|\varphi_n\|_V(x) & \text{if } x \in X_n \end{cases}$, so let $w \in \|\varphi\|_V(x)$ for $x \in X_i$, then $w \in \|\varphi_i\|_V(x)$ so by IH $(x, w) \in \|\overline{\varphi_i}\|_{\overline{V}}(x)$, and because $x \in X_i, (x, w) \in X_i, (x, w) \in \|p_i \wedge \overline{\varphi_i}\|_{\overline{V}}(x)$ and thus $(x, w) \in \|(p_1 \wedge \overline{\varphi_1}) \vee \dots \vee (p_n \wedge \overline{\varphi_n})\|_{\overline{V}} = \|\overline{\varphi}\|_{\overline{V}}$.

Now $(x, w) \in \|\overline{\varphi}\|_{\overline{V}} = \|(p_1 \wedge \overline{\varphi_1}) \vee \dots \vee (p_n \wedge \overline{\varphi_n})\|_{\overline{V}}$. Take i such that $(x, w) \in \|p_i \wedge \overline{\varphi_i}\|_{\overline{V}}$ then we have $x \in X_i$ and (by IH) $w \in \|\varphi_i\|_V(x)$, and by definition of $\|[\varphi_1, \dots, \varphi_n]\|$ then $w \in \|\varphi\|_V(x)$. \square

5 Conclusion and Future Work

In this report we have shown a coalgebraic representation of Büchi automata. The construction relies upon two key ideas: working in the Kleisli category for the monad \mathcal{P} and deriving two separate commuting diagrams for the accepting and non-accepting states and obtaining the right words by utilizing fixed point equations from these two mappings.

We explained the model in the Kleisli category in Section 3.1 by showing how to construct a final coalgebra for finite words for a nondeterministic system. Subsequently we constructed a weakly final coalgebra to additionally obtain the infinite words within such a system. Building upon these ideas we derived the coalgebraic construction for Büchi automata in Section 3.2, making use of the modal mu-calculus explained in Section 2.4.

We provided a proof for **lemma**, but not for Lemma 3.2, which is crucial for coincidence result in Theorem 3.3, and thus understanding why the construction indeed provides the words accepted by the Büchi automaton. Therefore, the first next step in the internship will be understanding the proof provided by [9].

After understanding the full proof of the coincidence result, we can start to think about replacing it using a different framework. Our goal is to replace it using a game semantics framework, which we briefly explained in Section 2.4 in relation to the modal mu-calculus. There, we showed how one can see the check whether a formula holds in a state as a two player game between a verifier and a refuter, who want to verify, respectively refute, that the formula holds. Our vision is that this view can be applied to whether a word is accepted by the coalgebraic model of a Büchi automaton, and that this could simplify the result.

Bibliography

- [1] André Arnold and Damian Niwinski. 2001. *Rudiments of mu-calculus*. Elsevier.
- [2] Steve Awodey. 2010. *Category theory*. OUP Oxford.
- [3] Erich Grädel, Wolfgang Thomas, and Thomas Wilke. 2003. *Automata, logics, and infinite games: a guide to current research*. Springer.
- [4] Ichiro Hasuo, Bart Jacobs, and Ana Sokolova. 2007. Generic trace semantics via coinduction. *Logical Methods in Computer Science* 3, (2007).
- [5] Bart Jacobs. 2004. Trace semantics for coalgebras. *Electronic Notes in Theoretical Computer Science* 106, (2004), 167–184.
- [6] Clemens Kupke and Yde Venema. 2008. Coalgebraic automata theory: basic results. *Logical Methods in Computer Science* 4, (2008).
- [7] John C Martin and others. 1991. *Introduction to Languages and the Theory of Computation*. McGraw-Hill NY.
- [8] Jan JMM Rutten. 2000. Universal coalgebra: a theory of systems. *Theoretical computer science* 249, 1 (2000), 3–80.
- [9] Natsuki Urabe, Shunsuke Shimizu, and Ichiro Hasuo. 2016. Coalgebraic Trace Semantics for Buechi and Parity Automata. In *27th International Conference on Concurrency Theory (CONCUR 2016)*, 2016. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 1–15.
- [10] Moshe Y. Vardi. 1996. An automata-theoretic approach to linear temporal logic. In *Logics for Concurrency: Structure versus Automata*, Faron Moller and Graham Birtwistle (eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 238–266. https://doi.org/10.1007/3-540-60915-6_6

A Proofs

Proof of Lemma 3.1: First we unfold some definitions:

$$(Jd)^{-1} = J(d^{-1}) \text{ and } d^{-1} = \text{cons and } J = \eta_{\Sigma^\omega}, \text{ so } J \circ d^{-1} = \eta_{\Sigma^\omega} \circ \text{cons.}$$

$\overline{F}[u_1, \dots, u_n] = \lambda_{\Sigma^\omega} \circ (\text{id} \times (u_1 + \dots + u_n))$ so let us call $u_1 + \dots + u_n = \text{beh}$ and see that $\text{id} \times \text{beh} : (\Sigma \times X) \rightarrow (\Sigma \times \mathcal{P}(\Sigma^\omega))$, maps a pair (σ, x) to $(\sigma, \text{beh}(x))$, i.e. σ and the language accepted by x . Combining with the natural transformation $\lambda : (\Sigma \times \mathcal{P}(\Sigma^\omega)) \rightarrow \mathcal{P}(\Sigma \times \Sigma^\omega)$ defined by $\lambda(\sigma, W) = \{(\sigma, w) \mid w \in W\}$ we get $\overline{F}[u_1, \dots, u_n](\sigma, x) = \{(\sigma, w) \mid w \in \text{beh}(x)\}$

$c_i = c \circ \kappa_i : X_i \rightarrow \mathcal{P}(\Sigma \times X)$ in terms of the automaton is defined as $c_i(x) = \{(\sigma, x') \mid x' \in X, \sigma \in \Sigma, x' \in \delta(x)(\sigma)\}$ for $x \in X_i$.

Combining these, and writing out the Kleisli composition in terms of functions in **Sets** we get:

$$(Jd)^{-1} \odot \overline{F}[u_1, \dots, u_n] \odot c_i = \mu_{\Sigma^\omega} \circ \mathcal{P}(\eta_{\Sigma^\omega} \circ \text{cons}) \circ (\mu_{\Sigma \times \Sigma^\omega} \circ \mathcal{P}(\lambda \circ (\text{id} \times (u_1 + \dots + u_n)))) \circ c_i. \quad (22)$$

Observing that $\mu_{\Sigma^\omega} \circ \mathcal{P}(\eta_{\Sigma^\omega} \circ \text{cons}) = \mathcal{P}(\text{cons})$, letting $u_1 + \dots + u_n = \text{beh}$ again and combining $\mathcal{P}(\lambda \circ (\text{id} \times \text{beh}))$ and c_1 by using our observations from above we obtain, for an $x \in X_i$:

$$\begin{aligned} & (\mu_{\Sigma^\omega} \circ \mathcal{P}(\eta_{\Sigma^\omega} \circ \text{cons}) \circ (\mu_{\Sigma \times \Sigma^\omega} \circ \mathcal{P}(\lambda \circ (\text{id} \times (u_1 + \dots + u_n)))) \circ c_1)(x) \\ &= \mathcal{P}(\text{cons})(\{(\sigma, w) \mid x' \in X, x' \in \delta(x)(\sigma), w \in [u_1, \dots, u_n](x')\}) \\ &= \{\sigma \cdot w \mid x' \in \delta(x)(\sigma), w \in \text{beh}(x')\} = \Diamond_\delta(\text{beh})(x) \end{aligned} \quad (23)$$

□